



Gradient-loss components vs reconstruction SSIM — D2 sweep (28 configs)

cos_sim vs SSIM — saturated, weakly predictive
(spread at cos~0.94: SSIM 0.18-0.55)



TV_final vs SSIM — strong inverse relation
(low TV final $\Rightarrow$ high SSIM)



total loss vs SSIM — best at left edge
(approx -0.95)

