WEXAC GPU Jupyter – Corrected Step-by-Step Guide (No $PORT Bug)

This version fixes the common SSH port-forwarding mistake. Environment variables (like **$PORT**) do NOT carry between SSH sessions. Always paste the numeric port explicitly.

1. **1. Login to WEXAC**
   ssh wexac
2. **2. Clear old jobs**
   bjobs -u $USER
   bkill 0
3. **3. Start interactive GPU job**
   bsub -q interactive-gpu -R "rusage[mem=8GB]" -gpu "num=1:j_exclusive=no:gmem=8GB" -Is bash
4. **4. Activate environment on GPU node**
   module load miniconda
   source /apps/easybd/programs/miniconda/24.11_environmentally/etc/profile.d/conda.sh
   conda activate /home/projects/galvardi/yoado/.conda/envs/rec
5. **5. Choose free port and start Jupyter (GPU node)**
   PORT=8899
   while lsof -nP -iTCP:$PORT -sTCP:LISTEN >/dev/null 2>&1; do PORT=$((PORT+1)); done
   echo "Using port $PORT"
   jupyter lab --no-browser --ip=127.0.0.1 --port=$PORT --ServerApp.port_retries=0
6. **6. Copy the printed port number**
   Example: **8899**
7. **7. Tunnel login2 → GPU (new login2 shell)**
   ssh -N -L 127.0.0.1:8899:127.0.0.1:8899 lgnXX
8. **8. Tunnel Mac → login2**
   ssh -N -L 8899:127.0.0.1:8899 wexac
9. **9. Open Jupyter**
   http://127.0.0.1:8899/lab?token=PASTE_TOKEN
10. **10. Restart kernel if needed**
    Kernel → Restart Kernel